

Project name: Next-Day Stock Price Forecast

Subject: Deep Learning

Project Link: <https://github.com/paolosilv/deep-learning2023>

Student: Paolo Silvestri

Student ID: 521343

Project Goals

Next day stock closing price (short-term prediction)

Should I buy or sell for tomorrow?

Data

Stock Symbol: AAPL -> APPLE

Provider: Yahoo Finance

Data Period: 2010-01-01 -> 2023-09-21

Primary Data Structure:

	Open	High	Low	Close	Adj. Close	Volume
Date						
2010-01-04	7.622500	7.660714	7.585000	7.643214	6.487534	493729600
2010-01-05	7.664296	7.699643	7.610071	7.656429	6.496751	601904800
2010-01-06	7.656429	7.696786	7.526786	7.534643	6.395379	552160000
2010-01-07	7.562500	7.571429	7.466071	7.520714	6.383556	477131200
2010-01-08	7.510714	7.571429	7.466429	7.570714	6.425996	447610800
...
2023-09-14	174.000000	176.100008	173.580002	175.740005	175.740005	60895800
2023-09-15	176.479996	176.500000	173.820007	175.009995	175.009995	109205100
2023-09-18	176.479996	179.380005	176.169998	177.970001	177.970001	67257600
2023-09-19	177.620004	179.630005	177.130005	179.070007	179.070007	51826900
2023-09-20	179.259995	179.699997	175.399994	175.490005	175.490005	58333200

Data Pre-Processing

Just the «Close» column feature

Data scaling -> MinMaxScaler

Training sequence -> 5 days

Data Splitting

Train size -> 85% -> 2929 rows

Test size -> 15% -> 518 rows

Batch size -> 64

Model Structure

LSTM -> 2 Layers | 64 Hidden Size

Fully Connected -> 1 Output

```
class StockForecastingLSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, output_size):
        super(StockForecastingLSTM, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)
        c0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)
        out, _ = self.lstm(x, (h0, c0))
        out = self.fc(out[:, -1, :])
        return out
```


Loss and Optimizer

Loss Metric -> MSE (Mean Squared Error)

$$= (1/n) * \sum (y_i - \hat{y}_i)^2$$

Optimizer -> ADAM (Adaptive Moment Estimation)

Training

Epochs -> 150

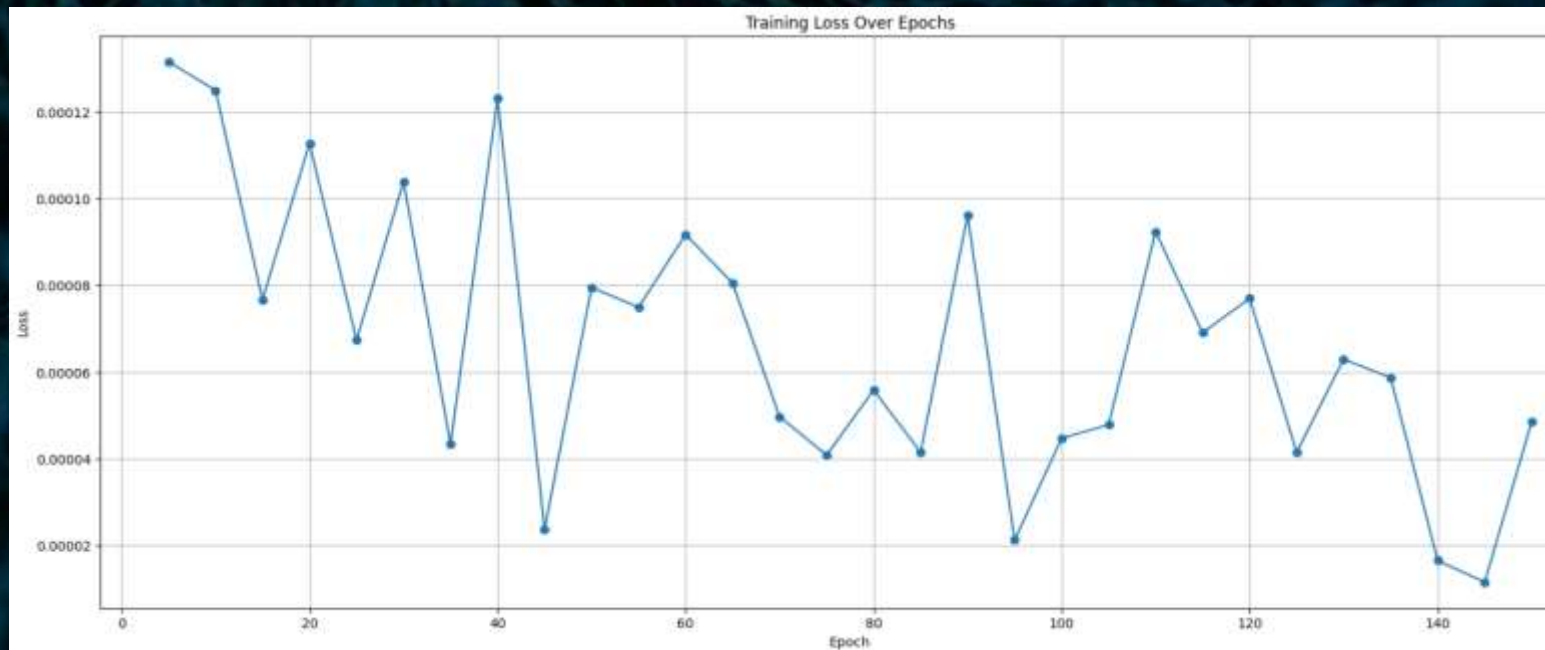
Learning Rate -> 0.001

```
# Training the model
losses = []
epochs = []
for epoch in range(num_epochs):
    for batch_seq, batch_target in train_loader:
        batch_seq = batch_seq.to(device)
        batch_target = batch_target.to(device)

        # Forward pass
        outputs = model(batch_seq)
        loss = criterion(outputs, batch_target)

        # Backward pass and optimization
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    if (epoch+1) % 5 == 0:
        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.6f}')
        epochs.append(epoch+1)
        losses.append(loss.item())
```



```
Epoch [5/150], Loss: 0.000132
Epoch [10/150], Loss: 0.000125
Epoch [15/150], Loss: 0.000077
Epoch [20/150], Loss: 0.000113
Epoch [25/150], Loss: 0.000067
Epoch [30/150], Loss: 0.000104
Epoch [35/150], Loss: 0.000043
Epoch [40/150], Loss: 0.000123
Epoch [45/150], Loss: 0.000024
Epoch [50/150], Loss: 0.000079
Epoch [55/150], Loss: 0.000075
Epoch [60/150], Loss: 0.000092
Epoch [65/150], Loss: 0.000081
Epoch [70/150], Loss: 0.000050
Epoch [75/150], Loss: 0.000041
Epoch [80/150], Loss: 0.000056
Epoch [85/150], Loss: 0.000042
Epoch [90/150], Loss: 0.000096
Epoch [95/150], Loss: 0.000021
Epoch [100/150], Loss: 0.000045
Epoch [105/150], Loss: 0.000048
Epoch [110/150], Loss: 0.000092
Epoch [115/150], Loss: 0.000069
Epoch [120/150], Loss: 0.000077
Epoch [125/150], Loss: 0.000042
Epoch [130/150], Loss: 0.000063
Epoch [135/150], Loss: 0.000059
Epoch [140/150], Loss: 0.000017
Epoch [145/150], Loss: 0.000012
Epoch [150/150], Loss: 0.000048
```

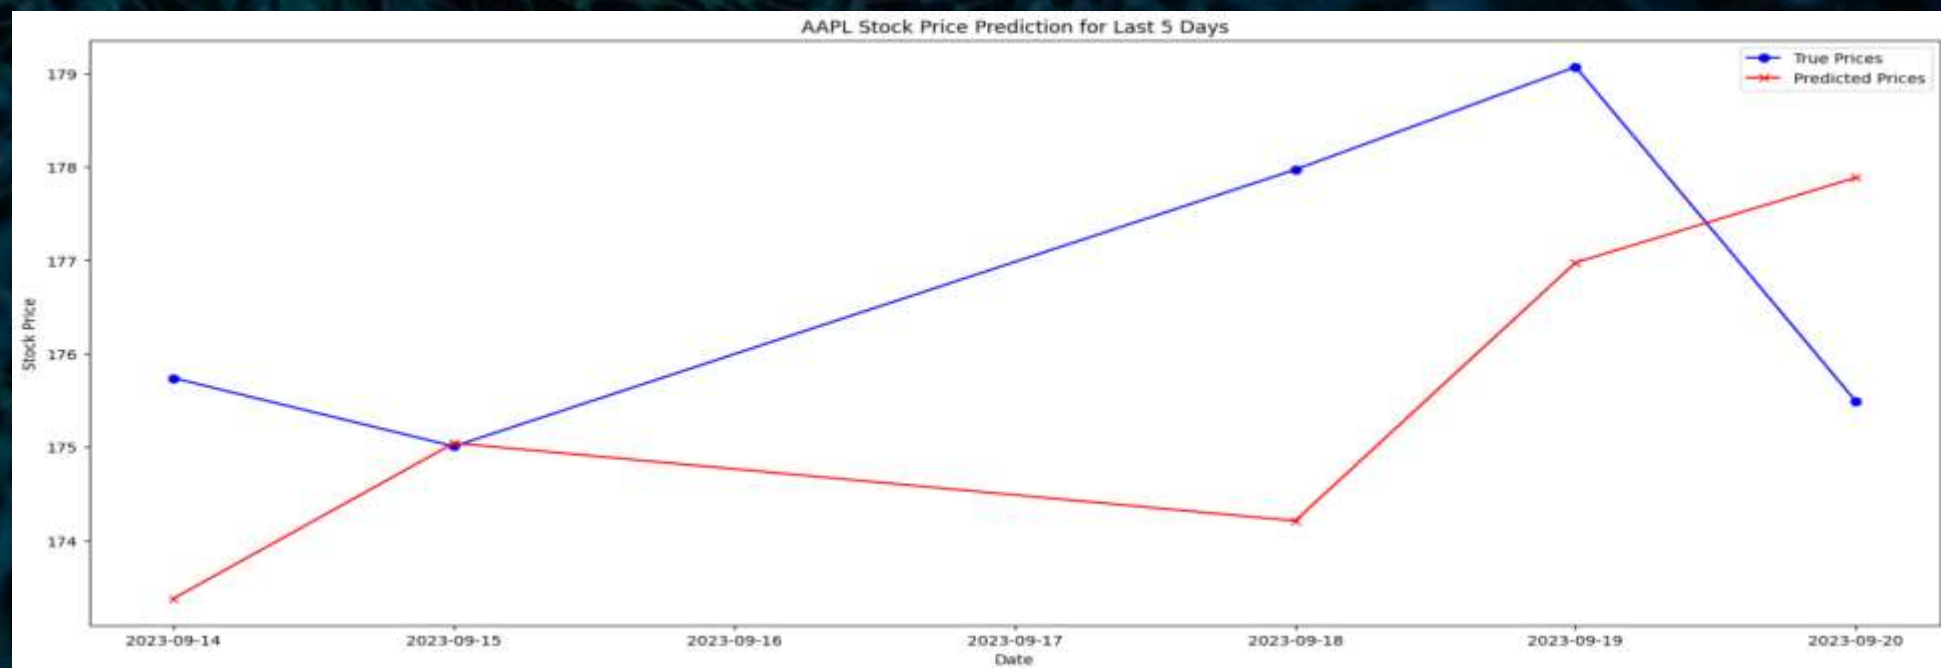
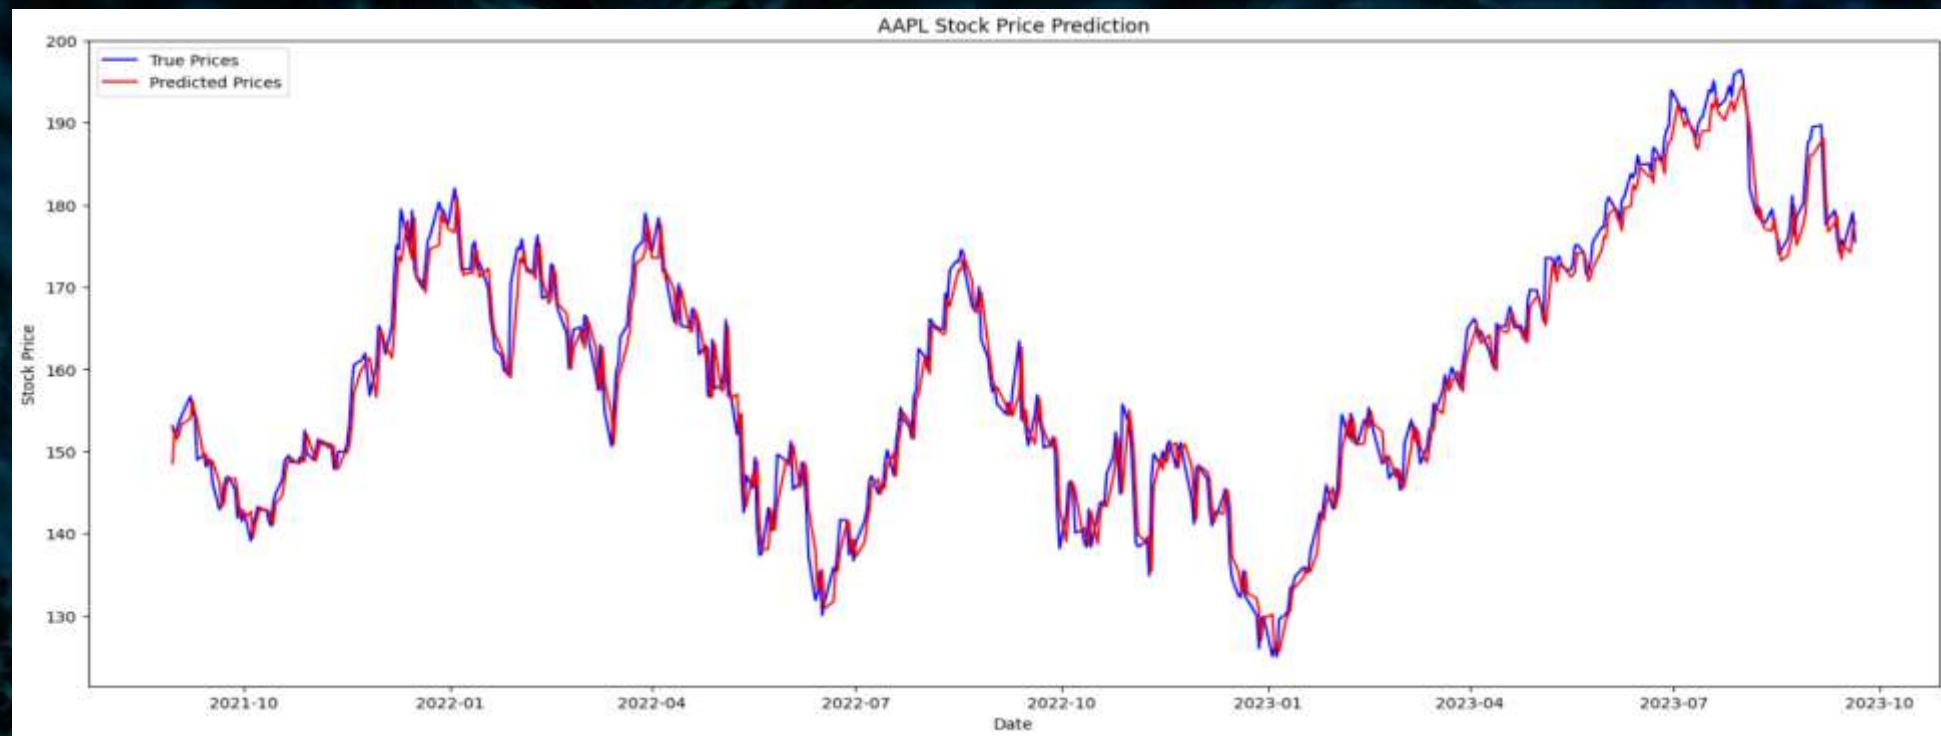

Testing

```
# Evaluate the model
model.eval()
test_seq = Variable(test_seq)
with torch.no_grad():
    predicted = model(test_seq).cpu().numpy()
    predicted = scaler.inverse_transform(predicted)
    true = scaler.inverse_transform(test_target.cpu().numpy())
```

MSE -> 8.7010

Mean Squared Error (MSE): 8.7010
Root Mean Squared Error (RMSE): 2.9497
Mean Absolute Error (MAE): 2.2641
R-squared (R²): 0.9661

Buy or Sell Accuracy -> 53.38%



Next Day Prediction -> 2023-09-21

Buy or Sell for tomorrow?





Thank You